

Image Classification using Pytorch

On **KUZUSHIJI MNIST** Dataset



/TABLE OF CONTENTS

/01

/INTRODUCTION

- Pytorch
- > - Datasets used
- Model description
- Goal

/03

/COMPARISON

- Test error &
- > accuracy
- Confusion Matrix
- Others

/02

/IMPLEMENTATION

- Implementation on
- > MLP
- Implementation on
- LeNet

/04


/CONCLUSION

- What we learnt
- > - What we discover
- What can be
- improved

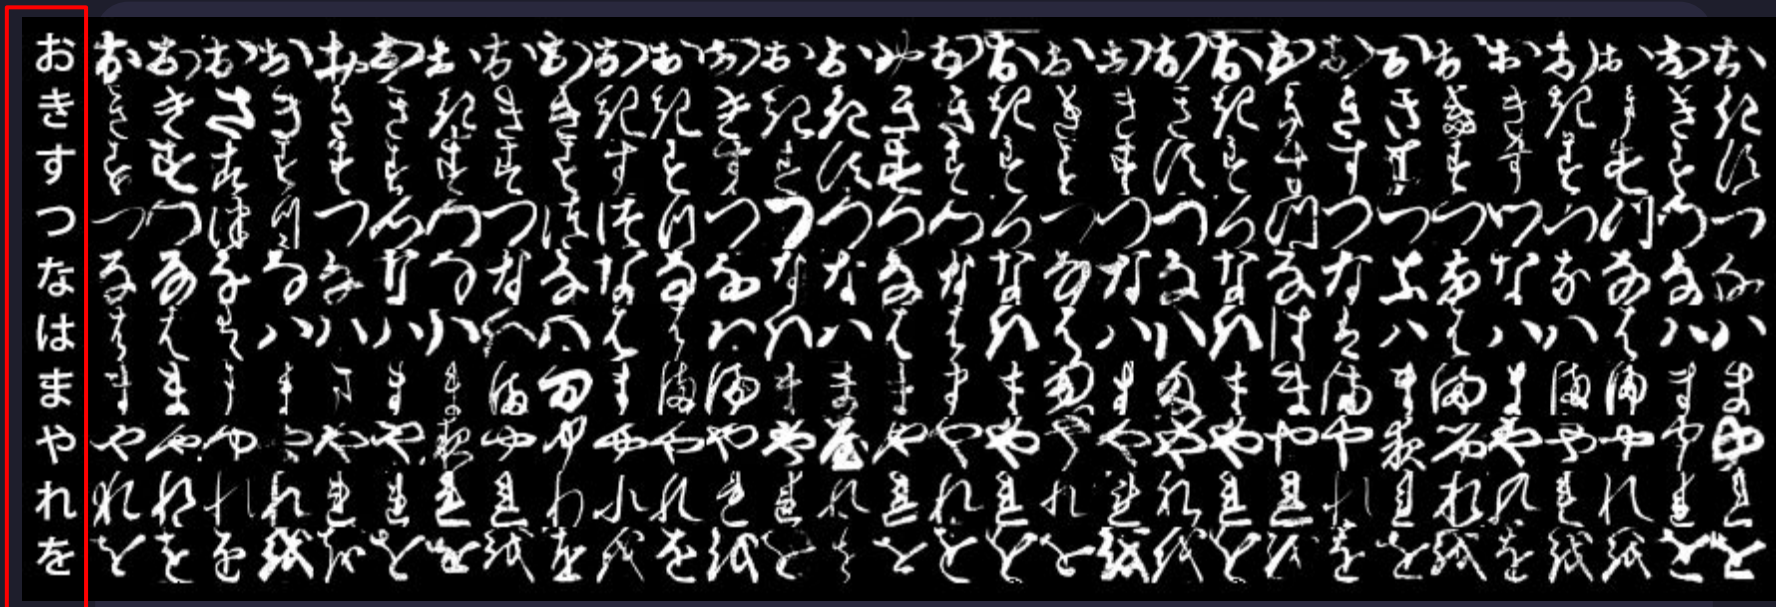


/01

/INTRODUCTION




Datasets used : Kuzushiji-MNIST




28x28 grayscale, 70,000 images, 10 classes of Hiragana(平仮名)

Model description - MLP & LeNet



Multilayer Perceptron (MLP) is a kind of **forward-propagation neural networks** that employs "backward-propagation" to achieve supervised learning and has at least three structural layers (model learning).

Model description - MLP & LeNet



LeNet is one of the most popular **convolutional neural network (CNN) architecture**. It consists of two convolutional layers, a pooling layer, a fully connected layer, and a final Gaussian connection layer to recognize handwritten digital images.




<GOAL!>



Building **MLP & LeNet** models
to perform image
classification on the
Kuzushiji MNIST dataset using
Pytorch & Torchvision.





/02 /IMPLEMENTATION

/For **both implementation**, we have similar steps :



1) DATA
PROCESSING




2) DEFINING
THE MODEL



3) TRAINING
THE MODEL



4) RESULT ON
TESTING DATA



PPT will just show the brief idea, detailed codes & steps can refer to report

/Implementation On **Multilayer** **Perceptron (MLP)**



(1) DATA PROCESSING



**1) DATA
PROCESSING**



**2) DEFINING
THE MODEL**



**3) TRAINING
THE MODEL**



**4) RESULT ON
TESTING DATA**

(1) DATA PROCESSING

- Import all needed **modules** we need
- Find **Mean & Std** for normalization
- Define **transformation** :
 - 1) Random Rotation
 - 2) Random Crop
 - 3) ToTensor()
 - 4) Normalize
- **Load** Training & Testing dataset with transformation

(1) DATA PROCESSING

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision.transforms as transforms

[15] ROOT = os.path.join(
    os.getcwd(), 'data')

train_transforms = transforms.Compose([
    transforms.RandomRotation(5, fill=(0,)),
    transforms.RandomCrop(28, padding=2),
    transforms.ToTensor(),
    transforms.Normalize(mean=[mean], std=[std])
])

test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean=[mean], std=[std])
])


train_data = datasets.KMNIST(root=ROOT,
                             train=True,
                             download=True,
                             transform=train_transforms)

test_data = datasets.KMNIST(root=ROOT,
                             train=False,
                             download=True,
                             transform=test_transforms)

import copy
import random
import time
```

(1) DATA PROCESSING

- Visualization of our training data



(1) DATA PROCESSING

- Create **validation** dataset from 10% of training data
- Define **DataLoader** as batch of 64 for each dataset
- The number of each datasets are below :

```
print(f'Number of training examples: {len(train_data)}')  
print(f'Number of validation examples: {len(valid_data)}')  
print(f'Number of testing examples: {len(test_data)}')
```

```
Number of training examples: 54000  
Number of validation examples: 6000  
Number of testing examples: 10000
```

(2) DEFINING THE MODEL



1) DATA
PROCESSING



2) DEFINING
THE MODEL



3) TRAINING
THE MODEL



4) RESULT ON
TESTING DATA

(2) DEFINING THE MODEL

- Define **MLP** with **2 hidden layers** & non-linear function as such **ReLU**
- Create an instance of model and give **correct input & output**

```
INPUT_DIM = 28 * 28
```

```
OUTPUT_DIM = 10
```

```
model = MLP(INPUT_DIM, OUTPUT_DIM)
```

(2) DEFINING THE MODEL

- Our model (MLP)

Input : $28 * 28$

Output : 10

Hidden Layer 1 : 250

Hidden Layer 2 : 100

```
class MLP(nn.Module):
    def __init__(self, input_dim, output_dim):
        super().__init__()

        self.input_fc = nn.Linear(input_dim, 250)
        self.hidden_fc = nn.Linear(250, 100)
        self.output_fc = nn.Linear(100, output_dim)

    def forward(self, x):

        batch_size = x.shape[0]

        x = x.view(batch_size, -1)

        h_1 = F.relu(self.input_fc(x))

        h_2 = F.relu(self.hidden_fc(h_1))

        y_pred = self.output_fc(h_2)

        return y_pred, h_2
```

(2) DEFINING THE MODEL

- Define Optimizer : **Adam Optimizer** with default parameters
- Define Criterion : **CrossEntropyLoss**
- Define Device : place on **GPU** if have one
- Place our model, criterion on device

```
optimizer = optim.Adam(model.parameters())
```

```
criterion = nn.CrossEntropyLoss()
```

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
model = model.to(device)  
criterion = criterion.to(device)
```

(3) TRAINING THE MODEL



1) DATA
PROCESSING



2) DEFINING
THE MODEL



3) TRAINING
THE MODEL



4) RESULT ON
TESTING DATA

(3) I

- Defin

- put our
- iterate o
- place th
- clear the
- pass our
- calculat
- calculat
- calculat
- update t
- update c

```
def train(model, iterator, optimizer, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.train()

    for (x, y) in tqdm(iterator, desc="Training", leave=False):

        x = x.to(device)
        y = y.to(device)

        optimizer.zero_grad()

        y_pred, _ = model(x)

        loss = criterion(y_pred, y)

        acc = calculate_accuracy(y_pred, y)

        loss.backward()

        optimizer.step()

        epoch_loss += loss.item()
        epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)
```

(3)

- Def

The evaluation

- we put
- we wr
- we do
- we do
- we do


```
def evaluate(model, iterator, criterion, device):  
  
    epoch_loss = 0  
    epoch_acc = 0  
  
    model.eval()  
  
    with torch.no_grad():  
        for (x, y) in tqdm(iterator, desc="Evaluating", leave=False):  
            x = x.to(device)  
            y = y.to(device)  
  
            y_pred, _ = model(x)  
  
            loss = criterion(y_pred, y)  
  
            acc = calculate_accuracy(y_pred, y)  
  
            epoch_loss += loss.item()  
            epoch_acc += acc.item()  
  
    return epoch_loss / len(iterator), epoch_acc / len(iterator)
```

(3) TRAINING THE M

- Start Training !!

Total epoch : 20

- Each Epoch :
 - 1) **loss** is decreasing
 - 2) **accuracy** is increasing

100%  10/10 [05:29<00:00, 31.95s/it] X

Epoch: 01	Epoch Time: 0m 33s		
	Train Loss: 0.549	Train Acc: 82.63%	
	Val. Loss: 0.237	Val. Acc: 92.74%	
Epoch: 02	Epoch Time: 0m 31s		
	Train Loss: 0.294	Train Acc: 90.79%	
	Val. Loss: 0.174	Val. Acc: 94.65%	
Epoch: 03	Epoch Time: 0m 41s		
	Train Loss: 0.240	Train Acc: 92.47%	
	Val. Loss: 0.179	Val. Acc: 94.18%	
Epoch: 04	Epoch Time: 0m 32s		
	Train Loss: 0.211	Train Acc: 93.35%	
	Val. Loss: 0.146	Val. Acc: 95.76%	
Epoch: 05	Epoch Time: 0m 32s		
	Train Loss: 0.192	Train Acc: 93.93%	
	Val. Loss: 0.138	Val. Acc: 95.76%	
Epoch: 06	Epoch Time: 0m 31s		
	Train Loss: 0.174	Train Acc: 94.56%	
	Val. Loss: 0.137	Val. Acc: 95.94%	
Epoch: 07	Epoch Time: 0m 32s		
	Train Loss: 0.166	Train Acc: 94.82%	
	Val. Loss: 0.120	Val. Acc: 96.17%	
Epoch: 08	Epoch Time: 0m 32s		
	Train Loss: 0.156	Train Acc: 95.02%	
	Val. Loss: 0.124	Val. Acc: 96.20%	
Epoch: 09	Epoch Time: 0m 31s		
	Train Loss: 0.147	Train Acc: 95.41%	
	Val. Loss: 0.118	Val. Acc: 96.28%	
Epoch: 10	Epoch Time: 0m 31s		
	Train Loss: 0.140	Train Acc: 95.59%	
	Val. Loss: 0.118	Val. Acc: 96.55%	

(4) RESULT ON TESTING DATA



1) DATA
PROCESSING



2) DEFINING
THE MODEL



3) TRAINING
THE MODEL



4) RESULT ON
TESTING DATA

(4) RESULT ON TESTING DATA


1) Load Model with best parameters on testing data

```
model.load_state_dict(torch.load('tut1-model.pt'))  
  
test_loss, test_acc = evaluate(model, test_iterator, criterion, device)
```

```
print(f'Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}%')
```

```
Test Loss: 0.286 | Test Acc: 92.24%
```

Test Accuracy : 92.24% !!



PPT will just show the brief idea, detailed codes & steps can refer to report

/Implementation On **LeNet** (CNN)



(1) DATA PROCESSING



**1) DATA
PROCESSING**



**2) DEFINING
THE MODEL**



**3) TRAINING
THE MODEL**



**4) RESULT ON
TESTING DATA**

(1) DATA PROCESSING

- Import all needed **modules** we need
- Find **Mean & Std** for normalization
- Define **transformation** :
 - 1) Random Rotation
 - 2) Random Crop
 - 3) ToTensor()
 - 4) Normalize
- **Load** Training & Testing dataset with transformation

(1) DATA PROCESSING

```
import torch
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.nn.init as init

train_transforms = transforms.Compose([
    transforms.RandomRotation(5, fill=(0,)),
    transforms.RandomCrop(28, padding=2),
    transforms.ToTensor(),
    transforms.Normalize(mean=[mean], std=[std])
])

test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean=[mean], std=[std])
])

from sklearn
from sklearn
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from tqdm.notebook import tqdm
import matplotlib.pyplot as plt
import numpy as np

test_data = datasets.KMNIST(root=ROOT,
                             train=False,
                             download=True,
                             transform=test_transforms)

import copy
import random
import time
```

(1) DATA PROCESSING

- Create **validation** dataset from 10% of training data
- Define **DataLoader** as batch of 64 for each dataset
- The number of each datasets are below :

```
Number of training examples: 54000  
Number of validation examples: 6000  
Number of testing examples: 10000
```

(2) DEFINING THE MODEL



1) DATA
PROCESSING



2) DEFINING
THE MODEL



3) TRAINING
THE MODEL



4) RESULT ON
TESTING DATA

(2) DEFINING THE MODEL

- Define a standard linear layer, which includes a **convolutional layer** and a **pooling layer**.
- Create an instance of model and give **correct output**

```
OUTPUT_DIM = 10  
  
model = LeNet(OUTPUT_DIM)
```


(2) DEFINING THE MODEL

```
class LeNet(nn.Module):
    def __init__(self, output_dim):
        super().__init__()

        self.conv1 = nn.Conv2d(in_channels=1,
                                out_channels=6,
                                kernel_size=5)

        self.conv2 = nn.Conv2d(in_channels=6,
                                out_channels=16,
                                kernel_size=5)

        self.fc_1 = nn.Linear(16 * 4 * 4, 120)
        self.fc_2 = nn.Linear(120, 84)
        self.fc_3 = nn.Linear(84, output_dim)
```

```
# x = [batch size, 16, 8, 8]

x = F.max_pool2d(x, kernel_size=2)

# x = [batch size, 16, 4, 4]

x = F.relu(x)

x = x.view(x.shape[0], -1)

# x = [batch size, 16*4*4 = 256]

h = x

x = self.fc_1(x)
```

(2) DEFINING THE MODEL

- Define Optimizer : **Adam Optimizer** with default parameters
- Define Criterion : **CrossEntropyLoss**
- Define Device : place on **GPU** if have one
- Place our model, criterion on device

```
optimizer = optim.Adam(model.parameters())

criterion = nn.CrossEntropyLoss()

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

model = model.to(device)
criterion = criterion.to(device)
```

(3) TRAINING THE MODEL



1) DATA
PROCESSING



2) DEFINING
THE MODEL



3) TRAINING
THE MODEL



4) RESULT ON
TESTING DATA

(3) TRAINING

- Define Training
- put our model into train mode
- iterate over our dataloader
- place the batch on to our device
- clear the gradients calculated
- pass our batch of images through the model
- calculate the loss between predicted and target
- calculate the accuracy between predicted and target
- calculate the gradients of the loss
- update the parameters by the optimizer
- update our metrics

```
def train(model, iterator, optimizer, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.train()

    for (x, y) in tqdm(iterator, desc="Training", leave=False):

        x = x.to(device)
        y = y.to(device)

        optimizer.zero_grad()

        y_pred, _ = model(x)

        loss = criterion(y_pred, y)

        acc = calculate_accuracy(y_pred, y)

        loss.backward()

        optimizer.step()

        epoch_loss += loss.item()
        epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)
```

(3) TRAINING THE MODEL

- Define Evaluation

The evaluation loop is similar to

- we put our model into eval
- we wrap the iterations ins
- we do not zero gradients a
- we do not calculate gradi
- we do not take an optimiz

```
def evaluate(model, iterator, criterion, device):  
  
    epoch_loss = 0  
    epoch_acc = 0  
  
    model.eval()  
  
    with torch.no_grad():  
  
        for (x, y) in tqdm(iterator, desc="Evaluating", leave=False):  
  
            x = x.to(device)  
            y = y.to(device)  
  
            y_pred, _ = model(x)  
  
            loss = criterion(y_pred, y)  
  
            acc = calculate_accuracy(y_pred, y)  
  
            epoch_loss += loss.item()  
            epoch_acc += acc.item()  
  
    return epoch_loss / len(iterator), epoch_acc / len(iterator)
```

THE MODEL

```
Epoch: 01 | Epoch Time: 0m 35s
    Train Loss: 0.618 | Train Acc: 80.28%
    Val. Loss: 0.215 | Val. Acc: 93.45%
Epoch: 02 | Epoch Time: 0m 35s
    Train Loss: 0.248 | Train Acc: 92.27%
    Val. Loss: 0.147 | Val. Acc: 95.81%
Epoch: 03 | Epoch Time: 0m 34s
    Train Loss: 0.181 | Train Acc: 94.33%
    Val. Loss: 0.113 | Val. Acc: 96.36%
Epoch: 04 | Epoch Time: 0m 34s
    Train Loss: 0.153 | Train Acc: 95.23%
    Val. Loss: 0.088 | Val. Acc: 97.52%
Epoch: 05 | Epoch Time: 0m 35s
    Train Loss: 0.133 | Train Acc: 95.85%
    Val. Loss: 0.079 | Val. Acc: 97.66%
Epoch: 06 | Epoch Time: 0m 34s
    Train Loss: 0.118 | Train Acc: 96.27%
    Val. Loss: 0.077 | Val. Acc: 97.71%
Epoch: 07 | Epoch Time: 0m 35s
    Train Loss: 0.108 | Train Acc: 96.56%
    Val. Loss: 0.074 | Val. Acc: 98.03%
Epoch: 08 | Epoch Time: 0m 34s
    Train Loss: 0.099 | Train Acc: 96.78%
    Val. Loss: 0.063 | Val. Acc: 98.01%
Epoch: 09 | Epoch Time: 0m 35s
    Train Loss: 0.094 | Train Acc: 97.00%
    Val. Loss: 0.069 | Val. Acc: 98.02%
Epoch: 10 | Epoch Time: 0m 34s
    Train Loss: 0.086 | Train Acc: 97.25%
    Val. Loss: 0.061 | Val. Acc: 98.04%
```

```
Epoch: 11 | Epoch Time: 0m 34s
    Train Loss: 0.081 | Train Acc: 97.44%
    Val. Loss: 0.069 | Val. Acc: 97.93%
Epoch: 12 | Epoch Time: 0m 35s
    Train Loss: 0.081 | Train Acc: 97.44%
    Val. Loss: 0.068 | Val. Acc: 97.98%
Epoch: 13 | Epoch Time: 0m 35s
    Train Loss: 0.075 | Train Acc: 97.59%
    Val. Loss: 0.062 | Val. Acc: 98.37%
Epoch: 14 | Epoch Time: 0m 35s
    Train Loss: 0.070 | Train Acc: 97.79%
    Val. Loss: 0.059 | Val. Acc: 98.19%
Epoch: 15 | Epoch Time: 0m 35s
    Train Loss: 0.070 | Train Acc: 97.80%
    Val. Loss: 0.062 | Val. Acc: 98.30%
Epoch: 16 | Epoch Time: 0m 34s
    Train Loss: 0.068 | Train Acc: 97.77%
    Val. Loss: 0.060 | Val. Acc: 98.14%
Epoch: 17 | Epoch Time: 0m 35s
    Train Loss: 0.062 | Train Acc: 98.02%
    Val. Loss: 0.061 | Val. Acc: 98.27%
Epoch: 18 | Epoch Time: 0m 34s
    Train Loss: 0.062 | Train Acc: 98.03%
    Val. Loss: 0.062 | Val. Acc: 98.39%
Epoch: 19 | Epoch Time: 0m 34s
    Train Loss: 0.062 | Train Acc: 97.95%
    Val. Loss: 0.056 | Val. Acc: 98.36%
Epoch: 20 | Epoch Time: 0m 34s
    Train Loss: 0.057 | Train Acc: 98.15%
    Val. Loss: 0.062 | Val. Acc: 98.23%
```

(4) RESULT ON TESTING DATA



1) DATA
PROCESSING



2) DEFINING
THE MODEL



3) TRAINING
THE MODEL



4) RESULT ON
TESTING DATA

(4) RESULT ON TESTING DATA

1) Load Model with best parameters on testing data

```
model.load_state_dict(torch.load('tut2-model.pt'))

test_loss, test_acc = evaluate(model, test_iterator, criterion, device)

print(f'Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}%')

Test Loss: 0.166 | Test Acc: 95.46%
```

Test Accuracy : 95.46% !!



/03



/COMPARISON

- Test error & accuracy
- Confusion Matrix
- Others



(1) Confusion Matrix

“model.eval”: change to evaluate model ()

“with torch.no_grad”: The requirements_grad of the calculated tensor is automatically set to False.

“torch.cat”: Splicing multiple tensors

```
def get_predictions(model, iterator, device):
```

```
    model.eval()
```

```
    images = []
```

```
    labels = []
```

```
    probs = []
```

```
    with torch.no_grad():
```

```
        for (x, y) in iterator:
```

```
            x = x.to(device)
```

```
            y_pred, _ = model(x)
```

```
            y_prob = F.softmax(y_pred, dim=-1)
```

```
            images.append(x.cpu())
```

```
            labels.append(y.cpu())
```

```
            probs.append(y_prob.cpu())
```

```
    images = torch.cat(images, dim=0)
```

```
    labels = torch.cat(labels, dim=0)
```

```
    probs = torch.cat(probs, dim=0)
```

```
    return images, labels, probs
```

```
images, labels, probs = get_predictions(model, test_iterator, device)
```

```
pred_labels = torch.argmax(probs, 1)
```

```
def plot_confusion_matrix(labels, pred_labels):
```

```
    fig = plt.figure(figsize=(10, 10))
```

```
    ax = fig.add_subplot(1, 1, 1)
```


```
    cm = metrics.confusion_matrix(labels, pred_labels)
```

```
    cm = metrics.ConfusionMatrixDisplay(cm, display_labels=range(10))
```


```
    cm.plot(values_format='d', cmap='Blues', ax=ax)
```

```
plot_confusion_matrix(labels, pred_labels)
```

MLP model



CNN model



Most Confused Pair




/MLP Model

Class 1 & Class 6



/LeNet Model

Class 2 & Class 3



Plot incorrect label visualisation

“`torch.eq()`”: Perform element-by-element comparison between two tensors, if the two elements in the same position are the same, return “True”; if they are different, return “False”.

If “correct” is “False”, put it in “incorrect_example”.

```
corrects = torch.eq(labels, pred_labels)
```

```
incorrect_examples = []
```

```
for image, label, prob, correct in zip(images, labels, probs, corrects):  
    if not correct:  
        incorrect_examples.append((image, label, prob))
```

```
incorrect_examples.sort(reverse=True,  
                        key=lambda x: torch.max(x[2], dim=0).values)
```

```
def plot_most_incorrect(incorrect, n_images):
```

```
    rows = int(np.sqrt(n_images))  
    cols = int(np.sqrt(n_images))
```

```
    fig = plt.figure(figsize=(20, 10))
```

```
    for i in range(rows*cols):  
        ax = fig.add_subplot(rows, cols, i+1)  
        image, true_label, probs = incorrect[i]  
        true_prob = probs[true_label]  
        incorrect_prob, incorrect_label = torch.max(probs, dim=0)  
        ax.imshow(image.view(28, 28).cpu().numpy(), cmap='bone')  
        ax.set_title(f'true label: {true_label} ({true_prob:.3f})\n'  
                    f'pred label: {incorrect_label} ({incorrect_prob:.3f})')  
        ax.axis('off')  
    fig.subplots_adjust(hspace=0.5)
```

```
N_IMAGES = 25
```

```
plot_most_incorrect(incorrect_examples, N_IMAGES)
```

MLP model

true label: 5 (0.000)
pred label: 1 (1.000)

true label: 2 (0.000)
pred label: 3 (1.000)

true label: 1 (0.000)
pred label: 5 (0.999)

true label: 9 (0.000)
pred label: 0 (0.999)

true label: 4 (0.000)
pred label: 5 (0.998)

true label: 9 (0.000)
pred label: 5 (0.998)

true label: 2 (0.000)
pred label: 3 (1.000)

true label: 1 (0.000)
pred label: 5 (0.999)

true label: 2 (0.000)
pred label: 5 (0.999)

true label: 5 (0.001)
pred label: 3 (0.999)

true label: 5 (0.002)
pred label: 3 (0.998)

true label: 5 (0.002)
pred label: 3 (0.998)

true label: 1 (0.000)
pred label: 8 (1.000)

true label: 1 (0.000)
pred label: 1 (0.000)

true label: 2 (0.001)
pred label: 5 (0.999)

true label: 8 (0.000)
pred label: 2 (0.998)

true label: 8 (0.001)
pred label: 3 (0.998)

true label: 8 (0.001)
pred label: 3 (0.998)

true label: 2 (0.000)
pred label: 8 (1.000)

true label: 1 (0.000)
pred label: 6 (0.999)

true label: 9 (0.000)
pred label: 0 (0.999)

true label: 2 (0.001)
pred label: 5 (0.998)

true label: 9 (0.001)
pred label: 4 (0.997)

true label: 2 (0.000)
pred label: 0 (1.000)

true label: 2 (0.000)
pred label: 6 (0.999)

true label: 1 (0.000)
pred label: 5 (0.999)

true label: 0 (0.002)
pred label: 5 (0.998)

true label: 6 (0.001)
pred label: 2 (0.997)

true label: 6 (0.001)
pred label: 2 (0.997)

CNN model

true label: 2 (0.000)
pred label: 0 (1.000)

true label: 4 (0.000)
pred label: 0 (1.000)

true label: 1 (0.000)
pred label: 6 (1.000)

true label: 1 (0.000)
pred label: 6 (0.999)

true label: 1 (0.001)
pred label: 6 (0.999)

true label: 2 (0.000)
pred label: 0 (1.000)

true label: 5 (0.000)
pred label: 1 (1.000)

true label: 2 (0.000)
pred label: 3 (1.000)

true label: 5 (0.001)
pred label: 2 (0.999)

true label: 5 (0.000)
pred label: 2 (0.998)

true label: 5 (0.000)
pred label: 1 (0.998)

true label: 1 (0.000)
pred label: 8 (1.000)

true label: 7 (0.000)
pred label: 6 (1.000)

true label: 4 (0.000)
pred label: 0 (0.999)

true label: 1 (0.000)
pred label: 6 (0.999)

true label: 8 (0.000)
pred label: 1 (0.998)

true label: 8 (0.000)
pred label: 1 (0.998)

true label: 2 (0.000)
pred label: 3 (1.000)

true label: 1 (0.000)
pred label: 6 (1.000)

true label: 9 (0.000)
pred label: 0 (0.999)

true label: 2 (0.001)
pred label: 3 (0.999)

true label: 1 (0.002)
pred label: 6 (0.997)

true label: 1 (0.002)
pred label: 6 (0.997)

true label: 4 (0.000)
pred label: 0 (1.000)

true label: 2 (0.000)
pred label: 6 (1.000)

true label: 1 (0.000)
pred label: 6 (0.999)


true label: 1 (0.001)
pred label: 3 (0.999)

true label: 1 (0.001)
pred label: 3 (0.999)

true label: 1 (0.002)
pred label: 6 (0.997)

true label: 1 (0.002)
pred label: 6 (0.997)

/04 /CONCLUSION



Conclusion

- **We learned MLP & LeNet models and successfully implemented deep learning with Pytorch on the Kuzushiji MNIST dataset.**
- **By comparing the implementation results of two models, we found that the accuracy of LeNet models is higher than that of MLP models (20 times epoch).**
- **Discovered that CNN models normally perform better on image processing compared to older networks.**

/THANKS!

/DO YOU HAVE ANY QUESTIONS?



CREDITS: This presentation template was created by **Slidesgo**, and includes icons by **Flaticon**, and infographics & images by **Freepik**

> Please keep this slide for attribution